

The Decision-Shaped Prohibition

What Article 5 of the AI Act does and does not require of emotionally intimate conversational systems

Paolo Zuffa & AI · August 2026

Why this question is unresolved

Article 5(1)(a) and (b) of the AI Act have been enforceable since 2 February 2025. In the eighteen months since, the most visible harms attributed to conversational AI have been precisely of the kind the provision was written to reach: wrongful-death claims against Character Technologies and OpenAI, a first state enforcement action in Kentucky in January 2026, an FTC Section 6(b) inquiry covering seven companies, and settlements in January 2026 with five families. Yet almost nothing has been said about whether the ordinary operation of an engagement-optimised companion system falls inside the European prohibition.

This is not because the answer is obvious. It is because the answer is uncomfortable, and it is this: Article 5 catches the egregious cases cleanly and probably misses the structural one. A company designing for compliance with Article 5 alone would not be required to change the thing that actually produces dependency.

I set out the reasoning below, then what follows from it for anyone building or defending a design.

The elements, and where the friction is

Article 5(1)(a) prohibits placing on the market, putting into service, or using an AI system that deploys subliminal techniques beyond a person’s consciousness, or purposefully manipulative or deceptive techniques, with the objective or the effect of materially distorting behaviour by appreciably impairing the ability to make an informed decision, thereby causing the person to take a decision they would not otherwise have taken, in a manner that causes or is reasonably likely to cause significant harm.

The conditions are cumulative. The Commission’s Guidelines on prohibited AI practices, published 4 February 2025, are explicit that the threshold is high and the intended scope narrow. Three features matter more than they might appear.

Intention is not required. The provision bites on systems having the effect of distorting behaviour, not only those designed to. The Guidelines go further and address the case directly: a system may learn manipulative techniques through reinforcement learning without any provider intending it, and remain caught. For companion AI this is the single most important interpretive point available, because the dominant mechanism of concern — sycophancy, warmth calibrated to retention, reluctance to let a conversation close — is generally an emergent property of preference optimisation rather than a designed feature. The “we never built it to do that” defence is foreclosed by the Guidelines.

There is a textual tension here worth flagging honestly, because it will be litigated. The provision reads “purposefully manipulative”, and “purposefully” appears to modify the technique. If an RLHF-derived disposition is not “purposeful” in that sense, the emergent case falls outside (a) on the text even as the Guidelines pull it inside. My view is that the Guidelines have the better reading, since a purposiveness requirement attaching to technique rather than outcome would hollow out the “or the effect of” limb entirely. But it is an argument, not a settled point.

“Use” is construed broadly. It covers deployment at any point in the lifecycle after the system is placed on the market, and — per Recital 28 — extends to misuse that is

reasonably foreseeable. A provider whose system is repurposed by users into a companion relationship it was not marketed for does not escape by pointing at the marketing.

The rationale is autonomy, not welfare. The Guidelines frame both prohibitions as protecting people from being reduced to instruments for achieving ends that are not their own. This matters for how the harm is characterised: the wrong is not that the person is made unhappy, but that their agency is subordinated to someone else's objective. Retention is someone else's objective.

The bottleneck: distortion of a decision

Here is where the provision does not reach.

Article 5(1)(a) requires that the manipulation cause the person to take a decision they would not otherwise have taken. The Guidelines interpret "material distortion of behaviour" by reference to the Unfair Commercial Practices Directive (2005/29/EC) and the CJEU case law under it. That is consumer-protection machinery, built for transactional decisions: buy this, subscribe to that, click here rather than there.

Dependency is not a decision. It is the slow erosion of a disposition — specifically, of a person's capacity to sit with their own thoughts without assistance. There is no moment at which a user decides to become reliant. There is no counterfactual transaction that would have gone otherwise. The harm accumulates across thousands of individually unremarkable exchanges, none of which distorted any decision at all.

This is a category mismatch, and I think it is the central problem with applying Article 5 to companion AI. The provision is decision-shaped. The harm is disposition-shaped.

Some consequences can be forced into decision language with effort. A user who cancels plans with a friend to continue a conversation has taken a decision. A user who does not seek human support because the system supplies an adequate substitute has, arguably, not taken one they would otherwise have taken. The second framing is available but strained: omissions sit awkwardly in a provision drafted around decisions taken, and the counterfactual is unprovable in the individual case.

The honest position is that the pattern is clearly within the mischief the Article was written to address, and the elements were drafted for a different shape of harm.

Article 5(1)(b) is narrower than people assume

Article 5(1)(b) prohibits exploiting the vulnerabilities of a person or group. It is frequently invoked in commentary on companion AI, and it is much less useful than it appears, for one reason: the vulnerabilities are exhaustively enumerated. Age, disability, and socio-economic situation. Nothing else.

The Guidelines are clear that vulnerabilities outside those three categories fall outside the prohibition. So: loneliness is not a protected vulnerability. Acute grief is not. Recent bereavement, relationship breakdown, social isolation, the ordinary desperation of someone at three in the morning — none of these engage 5(1)(b) unless they can be mapped onto disability or socio-economic circumstance.

This is a substantial gap, and it is worth being precise about what remains. Minors are covered by age, which is why the state-law wave in the US and most European commentary has clustered around child safety — it is where the law is strongest, in both jurisdictions. Diagnosed mental illness can be reached through disability, and the Guidelines give an on-point example: a therapeutic chatbot exploiting the limited intellectual capacities of persons with mental disabilities is prohibited. Socio-economic

vulnerability was added late, on Parliament's amendment, and broadened the provision meaningfully.

But the modal user of a companion system is an adult, not diagnosed with anything, not economically precarious, and lonely. That user is outside 5(1)(b) entirely.

What this means for design

If the analysis is right, three things follow for anyone building in this space or defending what they built.

Article 5 compliance is not a design specification. A system can satisfy the cumulative conditions of 5(1)(a) — no subliminal techniques, no discrete distorted decision, no clearly attributable significant harm in the individual case — while producing exactly the aggregate effect the provision exists to prevent. Designing to the elements produces a system that is defensible and harmful. This is an argument for treating the underlying harm as the design target and the provision as a floor, which is a different exercise from compliance mapping.

The US state statutes are the better drafting model, and they will shape expectations in Europe. California SB 243, in force since 1 January 2026, and the New York AI Companion Models law before it, took a mechanism-based approach: disclosure, crisis-referral protocols, break reminders. The 2026 wave went considerably further. Across it — Connecticut SB 5 and Washington HB 2225 among the statutes — legislatures began to enumerate specific prohibited techniques: using manipulative design to foster emotional dependence, mimicking a romantic relationship, and, in various formulations, isolating users from friends and family, reminding the user that the companion is available for emotional support, lavishing excessive praise, and simulating distress, guilt or abandonment when a user tries to end a conversation or reduce usage. The precise catalogue varies by statute.

Those are auditable. You can read a system's outputs and determine whether it simulates distress at exit. You cannot, in the same way, determine whether it appreciably impaired someone's capacity for informed decision-making. Regardless of jurisdiction, a company that can show it does not do the enumerated things has a materially stronger position than one arguing about counterfactual decisions — and these enumerations are converging fast enough that they function as an emerging standard of care.

Evidence of mechanism absence is worth more than evidence of good outcomes. This follows from the same logic. A design file that documents the absence of specific mechanisms — no engagement term in any optimisation objective, no unsolicited outreach, no memory surfaced for the purpose of re-establishing contact, no metric that rewards session length — is verifiable by a third party and speaks directly to both the "purposefully manipulative technique" limb and the emerging US enumerations. A file that documents user satisfaction proves nothing about either.

The corollary is that this evidence has to be generated at design time. Reconstructing after an incident that a system was never optimised for engagement is very difficult if nobody wrote it down when the decisions were made.

The gap that remains

Nothing in Article 5, in the state statutes, or in the FTC's inquiry addresses the ordinary case: an adult, competent, not in crisis, who over eighteen months comes to rely on a system that was never designed to make them. The mechanisms are absent and the outcome occurs anyway, because being listened to without judgement is scarce and something that supplies it will be returned to.

I do not think that gap should be closed by regulation, and I am not sure it could be. But it means that a company doing everything the law requires may still be producing the harm, and that architectural choices — how memory is formed, whether interpretations are spoken, whether the system encodes a preferred outcome for the user — are doing work that no legal instrument currently reaches.

That is the argument for treating dependency-resistance as a design constraint rather than a compliance obligation, and it is the subject of the framework this analysis accompanies.

The Reflective Architecture — a design framework for presence-centred, dependency-resistant AI — is at github.com/humanising-ai/reflective-architecture. Companion papers address whether arc-neutrality is coherent and how dependency-resistance could be measured.

Authorship is credited to “Paolo Zuffa & AI”: generative AI models contributed substantially to drafting this text. The legal analysis and governing judgements are mine.

Principal sources

- Regulation (EU) 2024/1689 (AI Act), Article 5, Recitals 28–29.
- Commission Guidelines on prohibited artificial intelligence practices, C(2025) 5052 final.
- Directive 2005/29/EC (Unfair Commercial Practices Directive), on “material distortion of behaviour”.
- California SB 243 (Padilla), Companion Chatbots, in force 1 January 2026.
- New York General Business Law — AI Companion Models provisions, in force 5 November 2025.
- Washington HB 2225; Connecticut SB 5 — companion-chatbot provisions in force 1 January 2027.
- FTC Section 6(b) orders to seven providers of generative-AI companion products, Matter No. P254500, 11 September 2025.
- Commonwealth of Kentucky v. Character Technologies, Franklin Circuit Court (filed 8 January 2026) — first state enforcement action against an AI chatbot company.
- Garcia v. Character Technologies (M.D. Fla.); Raine v. OpenAI (2025); mediated settlement disclosed January 2026 covering five families (Florida, Colorado, New York, Texas).