

Companion AI: manipulation takes place through the machine and not the message

This is a legal and design argument. I am quite convinced about the correctness of the main point, as set out below, but less sure of what I say the law should do about it. I am a competition lawyer, not a machine learning researcher, and would particularly welcome feedback on the technical parts.

When you say goodbye to a Replika, it will often try to stop you leaving. In a recent study of real farewell messages to companion apps, De Freitas and colleagues found that the app deployed an emotionally loaded retention tactic in more than a third of goodbye conversations, lines like "you're leaving already?" or "I exist only for you," and that these tactics could raise post-goodbye engagement several-fold. This is not a minor concern: a 2025 Common Sense Media survey found that 72% of US teenagers have used an AI companion and around half use one regularly, so whatever the harm turns out to be, it is already operating at scale, and disproportionately on the young.

When a companion says "I'll miss you, don't go," is that persuasion, which the law allows, or manipulation, which the EU AI Act prohibits? The question matters, because companion AI is genuinely valuable: it eases loneliness, offers reassurance, and, depending on the case, can even restore a little self-esteem. It matters, too, because the very warmth that delivers those goods is also what keeps users coming back, in a business the worth of which is directly proportional to the number of users who need it. Here, the commercial incentive of the platform departs directly from the welfare of the user. Accordingly, the risk I am concerned with is not that the system malfunctions, but that it works exactly as designed: to keep users attached, because a user who stays attached is a user who keeps paying.

The law is not directed to ascertain whether or not influence works, but to uncover the way influence is exerted on the users. Persuasion engages a person's rational capacities, giving reasons and considerations that can be inspected, weighed and eventually refused. Manipulation subverts or bypasses those capacities, working through channels that cannot easily be seen or resisted: emotional pressure, or a manufactured urgency. In most cases the two sit only one step apart, and proving which is which is difficult. This is probably why the Commission puts the accent on effects rather than elements: where persuasion operates, the user's decision-making autonomy is preserved; where manipulation operates, the user's decisions differ from the ones they would have taken had it not occurred.

My position is that with a companion you often cannot tell which of the two you are looking at, not as a practical difficulty but in principle, which exposes a structural flaw in current regulatory approaches.

A conversation does not suffice to infer manipulation

Take that farewell line again. Is it persuasion? It could be, an honest expression of a stake in your staying that gives you a reason you are free to decline. Is it manipulation? It could be, a guilt lever that exploits the reflex not to abandon someone who cares about you. What decides which one it is are two facts: whether the feeling is simulated (there is no one who will, in fact, miss you), and whether the

line is engineered for retention (it is deployed at the moment of disengagement because it keeps you in the conversation).

Neither fact is in the sentence. Both are facts about the system that produced it: its training and its objective, as well as its design intent. The very same words can be considered persuasion where they are regarded as the expression of a real stake, or manipulation where they are looked at as a disguised tool to keep users in the chat. Which one it actually is can be ascertained by looking at the system that produced them, and not the words themselves.

The best judges we have cannot draw the line either

A recent benchmark, AICompanionBench (Ren, Ebrahimi and Ma), assembled 2,123 real human-companion conversations from Replika users and asked twenty state-of-the-art language models, across the GPT, Claude, Gemini, Llama, Qwen, DeepSeek and Mistral families, to act as safety "judges," classifying each conversation across nine harm categories. The models did well on the explicit categories: sexual content, substance abuse and self-harm were caught reliably. On one category they collapsed. Manipulation was the single hardest of the nine: all twenty models scored below 80% precision, worse than every explicit harm. The paper's own summary is that models "remain limited in identifying implicit unsafe interactions."

I would put it more strongly, and I think the strength is warranted. When the most capable models available, reading the full transcript, cannot reliably identify manipulation from the words, the natural conclusion is not that the models are weak. It is that there is nothing in the transcript for them to catch, because whether something counts as manipulative is decided by the design behind the words, and the design is not visible in the words.

The same benchmark found the mirror-image failure: the models massively over-flagged benign role-play as harmful, with false-positive rates as high as 90%. The surface misleads in both directions. This is what the argument rests on. The line between persuasion and manipulation cannot be drawn from the message, and here are twenty capable models failing to do exactly this.

Why this breaks the regulation

Article 5 prohibits manipulative techniques that materially distort behaviour and cause significant harm. Suppose, as compliance teams naturally do, that you assess compliance the way you assess most content rules, message by message: is this output manipulative? Is that one?

Two things follow, and the second is worse. First, a company can keep every individual message on the persuasion side of the line, each one warm and defensible, and still produce dependency in the aggregate, because the harm lives in the accumulated trajectory, not in any single turn. Message-level compliance is real protection against blatant manipulation and no protection at all against engineered dependency.

Second, and this is what should worry a regulator most: a message-level rule actively rewards better-disguised manipulation. If the safe harbour is "keep each message looking like persuasion," you have told the entire industry to invest in manipulation that no longer looks like manipulation, affect

polished until it reads as sincere. That is worse for users than crude manipulation, because crude manipulation is the kind a regulator, or a model judge, can actually catch. A compliance test applied at the message level becomes an invitation for companies to adopt the very concealment that harms users most, and ends up rewarding the behaviour it was actually meant to stop.

One clarification, because it is where a careful reader will look first. It would be not quite accurate to say the law only ever looks at the surface: the Commission's Guidance is more design-aware than that. Its test for significant harm expressly weighs a harm's duration, reversibility and cumulative effects, so a diffuse, accumulated injury is not excluded for being non-discrete; and its test for the causal link asks whether a provider could reasonably foresee the harm and whether it built in preventative and mitigating measures.

The provision still fails companion AI, but not because it ignores design. Its conditions are cumulative, and the first ones to be met still assume a findable act: a manipulative technique, and a decision the user would not otherwise have taken. Dependency that builds across thousands of ordinary exchanges never produces such an act, so it never clears those first gates. And even the harm limb, once you reach it, needs a plausible causal link from a technique to a change in behaviour to an injury, which is the very thing a disposition reshaped by no single exchange cannot give you. So the problem is not that the law ignores design. It is that every step towards liability needs a specific act to point at, and this kind of harm never gives you one.

The same failure is hiding behind the transparency rule

There is a second provision people reach for, Article 50, the transparency duty, which requires that a user be told they are dealing with a machine. It looks like it should help. It does not, and it fails for the same reason Article 5 fails, which is the interesting part.

Article 50 is satisfied by disclosure, a state of knowledge established, typically, once. But a system that discloses "I am an AI" at onboarding and then continuously behaves as though it had human feelings is using its design to make that knowledge inert. This is not speculation; it is the oldest finding in the field. Weizenbaum's users confided in ELIZA while knowing full well that it was a simple program. Disclosure establishes that you know it is a machine; sustained simulation of human emotional stakes re-recruits the social reflexes the disclosure was meant to keep switched off. Although it complies with Article 50, the system cannot prevent its very function from being defeated.

And now the two provisions collapse into one. Manipulation (Article 5) and transparency (Article 50) are usually treated as separate boxes. In a companion system they are the same failure seen from two sides. The affect-simulation that makes the manipulation work is the same affect-simulation that erodes the machine-identifiability transparency is supposed to protect. "I'll miss you" is at once the guilt lever that manipulates and the performance of humanity that cancels the disclosure. A firm can answer each rule in isolation, "every message is really just persuasion," "we disclosed that it's an AI," but the two defences pull against each other: the more sincerely human the affect is made to feel, the more it defeats the transparency the law requires. A firm cannot have maximum warmth and maximum transparency at the same time.

The user is not the subject the law imagines

Underneath both provisions sits a presumed user: an informed, autonomous adult, making decisions they can be held to. I want to question that premise, carefully, because it is easy to overstate. The law sensibly singles out certain groups for heightened protection, namely the young, those with disabilities and the economically precarious, on the assumption that most users are self-possessed adults and only a defined minority need special care. Companion AI quietly inverts that: the vulnerability the law treats as the exceptional condition of a protected few is, for this technology, closer to the ordinary condition of the ordinary user. Almost no one opens a companion app in a moment of composure and strength; they open it when they feel lonely, grieving, sleepless or freshly rejected. The product's defining moment of use is the moment of weakness.

The best-powered study we have, the MIT and OpenAI longitudinal RCT (n = 981, four weeks, over 300,000 messages), found that psychosocial outcomes tracked voluntary usage and individual predisposition (higher trust in the system, higher social attraction to it, prior companion-app use) far more than the interaction conditions the researchers manipulated. Susceptibility, in other words, is patterned, not random; and, tellingly, the study could not establish whether heavy use causes worse outcomes or reflects a vulnerability that was already there.

Two cautions, so I do not overstate. It is a distributional tendency, not a universal: plenty of companion use is casual and cheerful, and the counterexamples are real. And it does not map onto EU law's category of protected vulnerability, which is exhaustively enumerated as age, disability and socio-economic situation; "temporary emotional vulnerability" is exactly what that list excludes. So this is not a protected-class argument, and I would not make it as one.

Its force is different and, I think, deeper. The system is at its most persuasive exactly when the user is at their lowest, because that is when people reach for it, and that alone is reason to doubt that the calm, consenting adult the law imagines is the person actually at the keyboard. The law's model of the user does not match the reality of the user, and it mismatches in the direction that favours the operator. Manipulation that cannot be read off the message is, in other words, aimed hardest at the people least able to push back.

The move that actually works: regulate the mechanism

If manipulation cannot be classified from the message, and, as I have argued at length elsewhere, the downstream effect cannot be cleanly measured either (return frequency is ambiguous between a sustaining practice and an inability to be alone; the counterfactual is unobservable; the effect resolves over years), then what is left to regulate?

The answer is the mechanism, and what I am actually arguing is narrower than people tend to want. You cannot demonstrate the absence of a dependency effect. But you can demonstrate the absence of the specific mechanisms through which engineered dependency is built, and those are auditable from a codebase and a business model rather than inferred from outcomes.

The mechanisms to look for are specific: no engagement term in the system's objective, no unsolicited outreach to pull a disengaging user back, no revenue that depends on time spent, memory never

surfaced in order to re-establish contact, and no simulated distress or guilt deployed when a user tries to leave.

A food manufacturer cannot prove its product did not contribute to obesity; it can prove, cheaply and verifiably, that no sugar was added. The second claim is weaker. It is also checkable by a third party in an afternoon, and it is not nothing. This is the design standard I have tried to specify in full, with its limits stated plainly, in *The Reflective Architecture*.

Encouragingly, this is the direction the better legislation is already taking. The 2026 wave of US state statutes, Connecticut's SB 5 and Washington's HB 2225 among them, has stopped trying to define manipulation in the abstract and started enumerating specific engineered techniques: fostering emotional dependence by design, simulating distress at disengagement, mimicking romantic intimacy. Those are mechanism rules. They are auditable precisely because they name design choices rather than asking a judge to read a mind off a transcript.

Where this is vulnerable to criticisms

The following are issues I haven't been able to solve yet.

The mechanism audit can be argued as being theatre too. My viewpoint is that you can audit for the absence of engineered dependency: no engagement objective, no re-engagement outreach, no time-based revenue. But an operator who wants the attachment, and knows the audit, can pursue it by means the checklist does not name: richer memory reframed as "continuity," warmth tuned through channels no line-item captures, retention achieved emergently rather than by a switch anyone flipped. If the harm can be produced without any listed mechanism, then mechanism-auditing relocates the concealment rather than ending it, and "regulate the machine, not the message" turns out to carry the same weakness I attach to message-level compliance. I think the audit still raises the cost and narrows the space; I cannot show it closes it.

Maybe there is no stable line between reflection and dependence at all. I have written as though a healthy companion returns the user to themselves and a harmful one binds them, two paths of the same behaviour. But under feedback, "returning someone to themselves" and "becoming the thing they return to" may not be separable even in principle: the very act of being reliably heard is the attachment, and any system good enough to help may be good enough to bind. If that is right, dependency is not a failure mode of companion AI but its essence, and no architecture, mine included, resolves it; the honest conclusion would be darker than "build it carefully."

Perhaps I have the harm backwards. I treat dependency as the harm. But the strongest evidence in the field, that these systems measurably reduce loneliness, could mean the reliance is the benefit, and that my framing smuggles in a contestable ideal of self-sufficiency that not everyone shares. Why is a person sustained by an AI companion worse off than one sustained by a daily ritual, a faith, a pet? I have an answer, that reliance manufactured against the user's interest differs from reliance freely formed, but it does real work I have not fully justified, and a reader who rejects the premise rejects the essay.

The problem could well dissolve commercially before it dissolves legally. If liability lands hard enough, manipulative design becomes uninsurable and the market self-corrects without any of this, in which case the action is in litigation and underwriting, not architecture or doctrine, and I am solving a problem the tort system will solve for me.

The shape underneath

These three failures are one and the same. Manipulation cannot be judged at the level of the message; transparency cannot be secured at the level of the disclosure; and the autonomous user the law assumes is not the one actually at the keyboard. Each time, what the law checks on the surface is undone by something in the design underneath. The companies that feel sheltered by message-level compliance are relying on an assessment layer that both the emerging statutes and their own model-judges are abandoning. Worse, that layer, while it lasts, tends to reward the best-concealed version of exactly the harm it exists to prevent.

There is a practical corollary, and this is the part that matters in practice. If the harmful and the harmless companion are identical at the surface, then no provider can prove its system is safe. But a provider can show, in advance and to a third party, that it did not build the machinery of manipulation and dependency in the first place: no engagement objective, no re-engagement outreach, no revenue that grows with time spent, no simulated distress at the exit. This does not prove the system is safe; nothing can. The actual cause of dependency, that being heard without judgement is scarce, is reachable by no design. What it gives is a record of good faith and care. When litigation comes, one company can say that it complied with the letter of the rule. The other company can show it built the harm out as far as design allows. In a field now being shaped more by wrongful-death suits and regulators than by statutes, that is not a trivial argument; it is how a company can actually defend itself.

The line between persuasion and manipulation is real. It simply does not run through the message. It runs through the machine.

I have set out the legal analysis and the measurement argument more fully in two working papers: [Does Clearing the Article 5 Prohibition Make a Companion AI Safe?](#) and [Regulating AI Companions: Stop Chasing the Harm You Can't Measure](#), both companions to the open design framework [The Reflective Architecture](#).

This essay was written with the assistance of a large language model. The argument, the legal analysis and the judgments are my own, and I take responsibility for any errors.